

IndiaFOSS 2026 Conference Programme Analysis

A structural assessment of the scheduled conference programme — what it covers well, where it thins out, and what a wholesome FOSS & digital-commons programme would still need.

This entire report was generated by an AI assistant (Claude, by Anthropic), not by IndiaFOSS or FOSS United. It is NOT an official FOSS United / IndiaFOSS publication, has not been reviewed or endorsed by the Programme Committee, and should be treated as an independent analytical starting point — not a final judgement on any session, speaker, or track. Only publicly available data was used to generate it — the published schedule, CFP proposals, and public FOSS United forum threads — nothing private, internal, or non-public. See the methodology note on the final pages for how classifications were made and where this analysis is most likely to be wrong.

138

SCHEDULED SESSIONS

69.6h

PROGRAMME MINUTES

10

TRACKS / DEVROOMS

3

CONFERENCE DAYS

Source: if2026-schedule.ods (General Track + 8 devrooms) plus the 8 pre-conference Workshop proposals from cfp-info.csv (approx. 2 hours each, held the day before the main conference), cross-referenced against the published CFP proposals at fossunited.org/c/indiafoss/2026/cfp, the FOSS United forum's devroom/BoF/panel call threads, and the public 75-booth listing at fossunited.org/indiafoss/2026/booths. This edition also adds a per-session "reach" figure — the seat capacity of the hall or room each session runs in — alongside the seven categories.md dimensions used throughout.

Executive Summary

IndiaFOSS 2026's programme is large, technically serious, and unusually community-minded. Counting the pre-conference workshop day, the conference runs 138 sessions across roughly 69.6 hours over three days, and builds its identity around four consistent pillars: a deep systems/embedded engineering cluster (Android internals, RTOS, compilers, open hardware, security), a broad FOSS community & civic-commons strand running through the General Track (contributor culture, GovTech, digital public goods, policy), a full dedicated Documentation & Technical Writing devroom — a track most conferences of this size skip entirely — and a pre-conference workshop day that gives eight topics a genuinely hands-on, 2-hour treatment. Average technical depth (5.6/10) and novelty (5.5/10) sit in a healthy "practitioner-intermediate" band: this is a conference of people telling war stories from real, shipped systems, not a research symposium and not a beginner bootcamp. That community-mindedness shows up in the numbers, too: 157 named speakers plus a 75-booth community floor put roughly 1 "producer" in the room for every 7–8 of the 2,401 registered attendees (Section 10) — an unusually dense participation rate next to most technology conferences. Sorting those 157 speakers by their own stated day job (Section 11) confirms it from another angle: startup/SME employees are the single largest professional background (25%), with non-profit/civil-society and students tied close behind (15% each) — though academia (6%) and government (1%) are both thin given how much research and GovTech content the programme itself covers.

This edition adds a new dimension on top of the usual categories.md schema: each session's **reach** — the seat capacity of the hall or room it's scheduled in. The finding is stark and structural: the 8 devrooms, despite carrying 54% of all sessions, are every one of them capped at a 120-seat room, while the General Track runs in Hall 1 (840 seats) or Hall 2 (250 seats). Weighting reach by duration, the General Track alone accounts for roughly 78% of the conference's total "seat-minutes" of exposure, despite being only 40% of sessions — see Section 12.

The programme is also lopsided in three predictable ways: within the two main conference days, format skews heavily toward short exposure (45% of all sessions are 15-minute-or-under lightning talks) at the expense of sustained depth; content skews toward the application layer even inside "systems" tracks; and the eight workshops, while a genuine strength, sit thematically apart from the conference's own technical core — none of them extend the RTOS, AOSP, Hardware, Compilers or Security devrooms into hands-on territory. Several other large, active areas of the FOSS/digital-commons world — data engineering & ML systems, accessibility, and full-stack open hardware beyond firmware — remain present only as a trace rather than as first-class content.

A note on who curates what

IndiaFOSS's programme has two different governance layers, and the gap analysis below is only fair if it's read with this split in mind. The General Track — including its Talks, BoFs and Panels, and the pre-conference Workshop day — is curated directly by IndiaFOSS's own Programme Committee (PC) from individual open proposals. The 8 devrooms (AOSP, Compilers, Documentation & Technical Writing, Cloud & DevOps, RTOS, Security, Design, Hardware) are each organised independently, who pick their own speakers and run their own internal review — but which devrooms exist at all was itself a PC decision: 12 devroom proposals were submitted to an open call, and the PC selected 8. BoF and Panel topics went through the same open-call-then-PC-selection model. "Independently organised" isn't uniform, though: only 3 of the 8 devrooms are run by an established, named community — Cloud & DevOps by Point Blank (who also handle a good chunk of the conference's on-ground volunteering), Compilers by LLVM Social Bangalore, and Hardware by absurd.industries — all three of whom are also listed as conference Community Partners (a separate category from sponsors — see Section 9). AOSP is a borderline case, led by a couple of individuals from the aosp-devs.org community. The remaining four — Documentation & Technical Writing, RTOS, Security, and Design — are run by individual volunteer organisers rather than a standing community or group. So a thin devroom (Hardware, at 6 sessions) is a signal about one organisation's bandwidth as much as community turnout — absurd.industries is simultaneously a Community Partner, this devroom's curator, and a booth exhibitor (Section 9), which is a lot of hats for one team. And a topic area with no devroom at all may mean a proposal for it existed and the PC didn't select it — which is exactly what happened for two of the gaps below. The tags mark which layer each gap sits in; where a specific rejected proposal is known, it's named.

Gap analysis — at a glance

INDEPENDENT DEVROOM Embedded systems depth —

AOSP + RTOS + Compilers + Hardware + Security together contribute 875 minutes of high-depth content (avg. depth 6.0–7.4/10), the technical core of the conference — run independently of the PC either way, though only Compilers and Hardware are backed by a standing named community (LLVM Social Bangalore, absurd.industries); RTOS and Security are individual-organiser devrooms and AOSP is a borderline case.

BOTH LAYERS FOSS community & contribution culture —

76 of 138 sessions touch "Community" as a domain; onboarding, maintainership, and governance are recurring threads in both the PC's General Track and the devrooms.

INDEPENDENT DEVROOM Documentation as a first-class

discipline — a full 11-session, 150-minute devroom, rare at this scale, and notably run by individual volunteer organisers rather than a standing community — a whole devroom's worth of quality without an established organisation behind it.

PROGRAMME COMMITTEE A genuine hands-on workshop

day — 8 pre-conference workshops (2 hours each, 960 minutes total) on Qt, OCaml, generative art/live-coding, AI agents, OpenStreetMap mapping, self-hosting, and digital security give real skill-building a dedicated day rather than a token slot.

PROGRAMME COMMITTEE **Civic tech / GovTech / digital-commons data** — judiciary mapping, GovTech archives, water intelligence, census/data-for-India work form a coherent public-interest thread inside the General Track.

INDEPENDENT DEVROOM **Compiler & PL depth** — LLVM, formal verification, Rust, Julia, WASM: a genuine deep-tech devroom, run by LLVM Social Bangalore, one of the conference's own Community Partners.

PROGRAMME COMMITTEE Workshops don't reach the

conference's own technical core — the 8 hands-on sessions (a Programme Committee-curated surface, same as the General Track) skew toward general-purpose software, creative-coding and civic themes; none is a hands-on extension of RTOS, AOSP, Hardware, Compilers or Security — the very devrooms with the highest average depth (6.0–7.4/10) get zero hands-on treatment, even though the PC clearly knows how to run a good workshop.

PROGRAMME COMMITTEE Data & AI systems have no

home track — and it's not for lack of proposals — a "Building with Open Data in India" devroom was proposed to the open call and not selected (8 of 12 proposals made it through); at the panel level, 2 of the 3 publicly-rejected panel proposals were AI-focused. What AI content survived (~20 sessions) sits scattered inside the PC's own General Track, Workshop day, Cloud & DevOps and Compilers. Partially offset by 5 data-focused community booths — see §9.

INDEPENDENT DEVROOM Open hardware beyond firmware

is the smallest devroom — only 6 sessions / 175 minutes, dwarfed by the software-adjacent embedded tracks (RTOS 10, AOSP 11). It's run by absurd.industries, simultaneously a conference Community Partner and a booth exhibitor (Section 9) — this reads less like weak community interest and more like one organisation carrying the whole devroom on top of its other conference commitments. Partially offset by 6 hardware-focused community booths — see §9.

PROGRAMME COMMITTEE Policy & legal depth is shallow

— a dedicated devroom for it was proposed and rejected — a "Tech Geopolitics and Policy" devroom was among the 4 of 12 proposals the PC didn't select. What's left is Legal/Policy sessions sitting inside the PC's own General Track or Workshop day, almost all as 15-minute lightning talks or BoFs. Partially offset by 4 policy/legal community booths — see §9.

WHO GETS THE MIC Reach is structurally capped for

devrooms — every devroom session runs in a 120-seat room regardless of depth or demand, while General Track sessions average a 411-seat reach; Hardware and Cloud & DevOps, this year's two most oversubscribed devrooms (Section 2 of the companion improvement plan), are also stuck at the same 120-seat ceiling as the least oversubscribed one (AOSP) — see Section 12.

WHO ATTENDS The conference itself is regionally

concentrated — by the dashboard's own regional breakdown, Karnataka and Kerala alone are 56% of registrations, and the five South Indian states together are 73% (Section 10). Not a programming-content gap, but a real limit on how "pan-India" the event's reach currently is.

Net assessment. The programme is wholesome for a systems-and-community-minded FOSS audience — a software developer, embedded engineer, or community organiser will find a full three days of relevant, well-paced content, including genuine hands-on time on the workshop day. It is less wholesome for a data/AI practitioner, a hardware maker beyond firmware, or a policy specialist wanting depth — and the one attendee still under-served on the skill-building front is exactly the person the deep-tech devrooms attract: someone who came for RTOS, AOSP, Hardware, Compilers or Security gets zero hands-on options in their own area, even though the workshop day proves the organisers know how to run one. Read against the two curation layers, essentially every gap here traces back to the Programme Committee's own selection decisions, not to community indifference. It's also worth reading the topical gaps alongside the reach gap: even where a devroom does cover a topic well, it structurally cannot put that content in front of as many people per session as a General Track talk can — a separate, amplification-level form of "thin," on top of whatever the topic coverage itself shows.

Detailed Overview

1. How the programme is built — and who curates it

The schedule has three layers, and critically, two different curators. A pre-conference Workshop day (8 sessions, 960 minutes) runs the day before the main event, each workshop given a full 2-hour hands-on slot. Then, across the two main conference days, the General Track (55 sessions, 1,875 minutes — 45% of all programme time) runs across three halls and rooms and is the conference's broadest surface: general FOSS talks, an invited talk, two panel discussions, eight BoFs, and one still-undetermined community-voted lightning slot ("Communi-Con", Section 12) sit alongside standard 30-minute talks. Both the General Track and the Workshop day are curated centrally by IndiaFOSS's own Programme Committee — every gap or strength found there is a direct reflection of the PC's own editorial choices. The remaining eight sheets are half-day mini-conferences ("devrooms") — AOSP, Compilers, Documentation & Technical Writing, Cloud & DevOps, RTOS, Security, Design, and Hardware — each concentrated into a single afternoon or morning block on one day, in a single room, and each organised independently, setting its own scope, review, and session count. That independence takes two different shapes, though: Cloud & DevOps, Compilers and Hardware are each run by an established, named community (Point Blank, LLVM Social Bangalore, and absurd.industries respectively — all three also listed as conference Community Partners, with Point Blank additionally handling much of the on-ground volunteering), AOSP is led by a couple of individuals from the aosp-devs.org community, and Documentation & Technical Writing, RTOS, Security and Design are each run by individual volunteer organisers rather than a standing group. Devrooms together contribute 75 sessions and 1,340 minutes (32%). This structure means a single attendee physically cannot sample more than a handful of devrooms across the two days, so the devroom line-up functions less like "an even split" and more like eight parallel mini-conferences that different sub-communities self-select into and self-organise — with the workshop day functioning as a distinct, opt-in fourth event, back under the PC's own curation.

2. Format and duration: breadth over depth on the main days

Of the 130 sessions across the two main conference days, 59 (45%) are lightning talks of 15 minutes or less, and 56 (43%) are standard 20–30 minute talks. 13 sessions run 60 minutes or longer — 12 are BoFs and panel discussions (open-floor formats, not deep single-speaker technical dives), and the 13th is the 75-minute Communi-Con slot (Section 12), the longest lightning-format booking on the entire programme despite having no content yet. There is effectively no 60-minute solo deep-technical talk slot anywhere in the two-day programme. Two devrooms — Cloud & DevOps (90% lightning talks) and RTOS (80%) — lean heavily toward rapid exposure over sustained depth (8–10 topics in roughly 2.5 hours), while Hardware and Design run zero lightning talks — every session there gets at least 20 minutes, and it shows in the narrative completeness of those proposals. This is a real editorial trade-off already made by the programme committee: devrooms with the most novel/frontier material (RTOS, avg. novelty 6.5/10) are also among the ones giving each idea the least time to land. The workshop day is the counterweight to all of this: its 8 sessions are uniformly 120 minutes, the longest, most sustained format anywhere in the conference.

3. What kind of work is this? Overall category and stack level

By minutes across all 138 sessions, the programme splits Software 2,370 min (57%) · FOSS Philosophy/Principles 1,030 min (25%) · Data 400 min (10%) · Hardware 300 min (7%) · not-yet-determined (Communi-Con) 75 min (2%). A quarter of scheduled time goes to community, governance, culture, policy and career topics rather than shipped technical artefacts — a signal that IndiaFOSS positions itself as a movement-building conference as much as a technical one. Within the technical stack-level tags, "Application" dominates at 65 tag-instances, ahead of Firmware (20), Frameworks/Middleware (19), Kernel (16), Operating System (14) and Libraries (14), and Hardware (11) — even the RTOS and AOSP devrooms, nominally "systems" tracks, spend much of their time at the firmware/application boundary (bootloaders, HALs, drivers) rather than deep kernel or silicon work. Genuine kernel- and silicon-level content (RISC-V crypto acceleration, TriCore bring-up, HackerFab semiconductor fabrication, sandbox kernels) exists but is a minority inside a minority, and the workshop day — being introductory by design — adds to the application-layer weight rather than the kernel/hardware end.

4. Domains served

Tagged against the domains in categories.md, sessions cluster heavily around Community (76) and Industry (65), with Research (30), Education/Academia (20), Government (14), Not-for-Profit Sector (12) and Legal/Policy (9) trailing well

behind (sessions can carry more than one domain tag). Government and civic-tech content is real and often high-quality — the Telangana Oracle-to-PostgreSQL migration, the NHAI AI-workflow talk, the Judiciary Entity Map, the CBSE exam-security disclosure story — but at 14 tagged sessions it is a minor thread relative to the scale of India's GovTech/DPI ecosystem that FOSS United's own General-Track panel ("From DPI to People's Digital Infrastructure") explicitly calls out as underserved. Education/Academia content gets a boost from the workshop day (Qt, OCaml and generative-art workshops are explicitly pitched at learners) but is still sparse given how much of the FOSS pipeline (GSoC, campus communities, COSH-style centres) depends on it.

5. Depth vs. novelty, and who the programme is for

Across the 137 sessions with a determinable score (Communi-Con excluded), average depth is 5.6/10 and average novelty 5.5/10, with RTOS (7.4 depth) and Security (6.8 depth) the deepest tracks, and the Workshop day (4.8 depth, by design — these are onboarding sessions for newcomers) and Documentation & Technical Writing (3.7 depth) intentionally the most accessible. The dashboard below (Figure 1) breaks depth vs. novelty out by overall category: Hardware sessions cluster at high depth/high novelty (little hand-holding, aimed at practitioners already in the space), while FOSS Philosophy/Principles sessions cluster at low-to-mid depth (accessible almost by design, since the goal is persuasion and culture-building rather than technique transfer). Audience tags mirror this: software developers are the explicit or implicit audience for 99 of 138 sessions, with hackers (39), researchers (33) and business people (23) next; educators/professors (10) and designers (10) are among the least served named audiences from the categories.md list, and venture capitalists are the one named audience with zero sessions.




Figure 1. Programme composition across six metrics tracked in this assessment: session distribution by track, format mix, overall-category time split, domain coverage, the depth/novelty relationship by category (bubble size = session duration), and — new this edition — reach, the number of sessions running in each room-capacity tier. Track labels in the top-left panel are tagged (PC) for the Programme Committee-curated surface (General Track + Workshops) and (dev) for the 8 independently-curated devrooms.

6. The ambient-AI observation

No track or devroom is named for AI or machine learning, yet AI/agent-related themes surface across a meaningful share of the programme by scheduled minutes — sandboxing AI agents, AI on-call agents, KV-cache systems work, AI-driven language assessment, AI evaluation frameworks for non-profits, one full hands-on workshop on building AI agents, and multiple policy sessions on AI-generated imagery and open-weight model harms. This is arguably the right editorial call for a conference about Free and Open Source Software rather than AI specifically — but it also means attendees looking for a coherent "open/local AI systems" thread have to assemble it themselves from talks and one workshop scattered across five unrelated tracks.

7. The pre-conference workshops, in detail

The 8 workshops are: Getting Started with Qt (desktop GUI development), Making Music with strudel.cc (live-coding, complete-beginner friendly), Digital Security for Journalists, Activists, and Everyday Internet Users (privacy/threat-modelling), Fun and Profit with OCaml (functional programming, run by IIT Madras's FP Launchpad), Art – meet Math – meet Code (generative art with p5.js), Building AI Agents End-to-End (agent frameworks, MCP, observability), Adding Places to OpenStreetMap (an outdoor mapping session), and Self-Hosting 101 (homelab/server setup). Read as a set, this is a genuinely well-rounded introductory slate — it spans desktop software, functional programming, civic mapping, digital rights, creative coding, and modern AI tooling, and three of the eight (strudel.cc, Art/Math/Code, OpenStreetMap) are explicitly designed for participants with no prior programming background, which is rare generosity in a conference this technical. The gap is thematic rather than qualitative: an attendee who wants to walk in the door already excited about RTOS, embedded Android, compilers, open hardware or security has no equivalent hands-on on-ramp — every one of those devrooms remains talk-only, even though the organisers have clearly demonstrated, via the other eight workshops, that they can run a hands-on session well. All 8 workshops also share the smallest reach of any session type on the programme — 35–40 seats each, Section 12 — a direct, mechanical consequence of hands-on formats needing small group sizes, not an editorial slight.

8. What didn't make it: the proposals the Programme Committee turned down

IndiaFOSS runs its devrooms, BoFs and panels through open public calls on the FOSS United forum, and the PC's own selection threads are public — which means this report can name specific rejected proposals rather than just inferring gaps from what's missing. Devrooms: 12 proposals were submitted to the March–May 2026 call-for-devrooms; the PC selected 8 (AOSP, Compilers, Documentation & Technical Writing, Cloud & DevOps, RTOS, Security, Design, Hardware) and turned down 4 — Tech Geopolitics and Policy (digital sovereignty, trade policy, regional FOSS adoption), Building with Open Data in India (climate-tech and civic-tech data quality), Open Cellular Communication (open-source 5G/Open RAN telecom), and FOSS at the Edge of Experience (creative coding, games, spatial computing). Panels: at least 5 were proposed for 2 slots; of the 3 unsuccessful proposals visible on the forum, 2 were explicitly AI-focused — Building a Global AI Commons and Who Pays for Cognition? Open Source in the Age of Organization-Hosted AI Agents — alongside Community Infrastructure Matters: Scaling People Like We Scale Platforms. BoFs: the organiser-supplied CFP submission record shows 14 formal BoF submissions this year — 11 accepted, only 3 rejected (Operationalising Transparency in State AI Deployment, AI for Security / Security for AI / Security from AI, and Organizing OpenStreetMap Mapping Parties) — a ~79% acceptance rate. That's a much higher hit rate than devrooms or panels: BoFs are the one format IndiaFOSS isn't short of room for.

Read together, this changes the shape of the gap analysis on page 2: the Data, Policy, Telecom and Creative-Experience gaps aren't blind spots nobody in India's FOSS community cares about — they're topics with organised advocates who put together a full devroom pitch and lost out in a competitive selection process, most likely on capacity grounds (8 of 12 is already a lot of parallel programming for a two-day conference). The AI panel proposals tell a similar story at smaller scale. None of this is a criticism of any single decision — a PC has to say no to good proposals every year — but it does mean the honest gap-analysis framing is "proposed and not selected," not "absent from the community." The companion report, [if2027-improvement-plan.pdf](#), turns this rejected-proposal record into concrete 2027 recommendations.

9. Beyond the talks: the booth floor and sponsors

Everything above is about the talk programme. IndiaFOSS also runs a 75-booth community floor (fossunited.org/indiafoss/2026/booths) — tables staffed by two representatives each, running one or both days, where projects and communities demo work and talk to attendees directly rather than present from a stage. The public listing groups them as FOSS Project (27), Community (29), Hardware (6), Student-led community/project (6), and a thematic "All Things Open" cluster (7); 55 of the 75 run both days.

The booth floor materially offsets three of the session-level gaps above — worth weighing before concluding a topic is genuinely absent from the conference. Open hardware gets real floor presence beyond the small devroom: absurd.industries, Ampere Works, PCB Cupid, Pling! (an open-source keyboard), T-01 The Explorer (robotics), and GOSH (Global Open Science Hardware) — six hardware-focused tables, roughly matching the devroom's own session count. Policy & legal, shut out at both the devroom and panel level, still shows up as SFLC.in (software freedom legal advocacy), Constitutional Observer, Takshashila Institution, and DAKSH & the Judicial Data Collaborative — four booths doing exactly the sustained policy work the talk programme couldn't fit in. Data & civic-data — the rejected devroom's actual subject — gets Data For India, Open Data Kerala, Samaajdata, OpenCity, and API Dash. What the booth floor does not reach: nothing in the list is disability-accessibility-focused (a few booths touch language localisation instead — Indic MediaWiki Developers, Sanchaya/ServantsOfKnowledge, Wikimedians of Kerala — a related but distinct gap), and nothing revisits the rejected Telecom or Creative-Experience devroom pitches beyond one mesh-networking booth (Janastu Anthillhacks).

Separately, IndiaFOSS lists 10 sponsors (Samagata Foundation, Zerodha, Frappe, Bruno, Ente, Red Hat, Tangled, Juspay Hyperswitch, Mecha Systems, Grafana) and, as a distinct category, 3 named community partners (Point Blank, absurd.industries, LLVM Social Bangalore) on the conference's main page — logos and hyperlinks only, no talk slot attached, so they sit outside this report's session-by-session classification entirely. They're a rough proxy for who's funding and helping run the event rather than what it covers: several sponsors (Frappe, Grafana, Red Hat, Juspay) map straight onto the programme's already-strong Cloud/DevOps and enterprise-software content, and two organisations (Mecha Systems as a speaker/booth presence, absurd.industries as a community partner) also appear elsewhere in the sections above — sponsorship, community partnership and content aren't fully separate worlds here, but neither list meaningfully shifts any of the gap findings.

10. Who's in the room: attendee stats and the producer-to-attendee ratio

IndiaFOSS's own public dashboard (fossunited.org/indiafoss/2026/stats) reports 2,401 total registrations as of this report's snapshot. The five paid ticket tiers it breaks out — Students (751), Enthusiast (478), Late Bird Enthusiast (230), Regular (144), and Contributor (38) — sum to exactly 1,641, precisely 760 short of that headline total. That gap is no longer a mystery: IndiaFOSS reserves 760 free tickets for speakers, volunteers, sponsors, booth staff, and a separately-run diversity scholarship, and the dashboard's ticket-tier breakdown covers only the 1,641 purchased tickets, not these 760 reserved passes — confirmed directly rather than inferred, unlike in this report's previous edition. Gender skews heavily male (1,208 / 75% Man, 367 / 23% Woman, 26 / 2% prefer not to say, 14 / 1% Non-binary, percentages as published), and geography skews heavily South Indian: Karnataka (604, 32%) and Kerala (458, 24%) alone are 56% of registrations by the dashboard's own regional breakdown, and Karnataka + Kerala + Tamil Nadu (198, 10%) + Andhra Pradesh (73, 4%) + Telangana (61, 3%) together account for 73% — this reads much more as a Bangalore-region flagship event than an evenly-distributed pan-India one, worth naming plainly alongside the topical gaps when asking "wholesome for whom." (Gender and region are optional profile fields, so their published percentages are computed against whatever subset of registrants answered them, not against the full 2,401 or the 1,641 ticket-purchaser figure — the dashboard doesn't state that base directly, so this report quotes its own published percentages rather than recomputing a different one.)

Speakers and community booths are usefully thought of as the conference's "producers", with everyone else as attendees. This report can now confirm what the previous edition could only guess at: producers draw from the same 760-ticket reserved pool as volunteers, sponsors and diversity scholars, and sit *inside* the 2,401 registration total rather than on top of it — resolving the ambiguity the last edition flagged about whether this ratio was additive or a genuine share of the room. Pulling the speaker field from all 138 session pages fetched for this analysis gives 157 unique named speakers across the talk programme (170 speaker-slots total, since 13 people speak twice). Add the 75-booth floor at two representatives each (150 more producer seats, per how each booth is staffed) and the upper-bound producer count is ~307 — well within the 760-ticket reserved pool, with the remaining ~453 reserved tickets accounted for by volunteers,

sponsor teams and diversity scholars this report doesn't count as "producers." Against the full 2,401-registration total, that's roughly 1 producer for every 7–8 attendees, or about 13% of the room — now a confirmed share-of-the-room figure, not an upper-bound guess. It's still not an exact headcount: there's no way to check whether a booth's two named representatives are also a session speaker, and 150 assumes every booth keeps the same two people the whole time rather than rotating volunteers. Even so, this is a strikingly dense participation rate for a conference this size — most technology conferences run closer to one speaker for every 30–50+ attendees.

11. Who's speaking: the producer community's own demographic mix

Attendee stats answer "who's in the room"; this section asks the same question of the 157 unique speakers themselves — not by gender or region (name-based guessing at either would be unreliable and isn't attempted here), but by professional background, since that's what the request behind this report specifically asked for. Every one of the 138 cached session pages also carries each speaker's own designation-and-organisation line (e.g. "Senior Software Engineer | Qt India Technology"); this report read all 157 of those lines by hand and sorted each speaker into one bucket. This is a single-analyst judgement call from a one-line self-description, not a verified census, and career stage and employer type inevitably overlap for some people (a "Founding Engineer" could as easily be filed under startup employee as founder — this report's rule was to reserve the founder/entrepreneur bucket for titles that explicitly say "Founder" or "Co-Founder"); treat the picture as directional.




Figure 2. All 157 unique speakers, sorted into one professional-background bucket each, from the designation and organisation each gave on their own proposal page.

Startup/SME employee is the single largest bucket (40, 25%) — more speakers work at a smaller product or services company than in any other category. Non-profit/civil-society (24, 15%) and Student (23, 15%) are close behind, and Large/Established Company adds another 21 (13%) — combine the three industry-adjacent buckets (startup, large company, founder) and industry practitioners are a 49% plurality, so this is a conference where most producers have a day job in industry even as its community-first identity remains real and visible. Founder/Entrepreneur is 16 speakers (10%), and speakers who described themselves too thinly to place with confidence ("Software Engineer" with no organisation, or a bare role like "Collaborator") land in Not stated/Unclear at 15 (10%).

Two buckets are genuinely thin. Academic/researcher speakers are only 6% (9 of 157) — light for a programme that includes a real research thread (PeerDAS, formal verification, genomics search, social-ecological data structures) — most of that content is delivered by industry engineers and PhD students rather than faculty. Government is the thinnest bucket of all, at just 1% (2 of 157) — both from NHA1 — despite Government being a tagged domain on 14 sessions elsewhere in this programme (Section 4): most of IndiaFOSS's GovTech and civic-infrastructure content is told by vendors, civic-tech builders and researchers about government systems, not by government employees themselves. Independent/freelance speakers are similarly rare (7, 4%) — a mild surprise for a movement historically associated with the solo hacker, though it may just mean solo contributors describe themselves by their day job instead.

12. Reach: which sessions get the biggest room

This edition adds a new per-session figure — **reach**, the number of seats in the hall or room each session is scheduled into. IndiaFOSS 2026 runs sessions in five room sizes: Hall 1 (840 seats), Hall 2 (250), Hall 3 (120), Room 1 (120), and Room 2 & 3 (35 seats each); the pre-conference Workshop day runs in its own dedicated Workshop Venue, capacity 35–40. Cross-referencing reach against every other dimension in this report surfaces a structural pattern the topical gap analysis alone can't see.

Room / capacity	Sessions	Who runs there
Hall 1 — 840 seats	19	General Track only
Hall 2 — 250 seats	25	General Track only
Hall 3 / Room 1 — 120 seats	75	All 8 devrooms, every session, no exceptions
Room 2 & 3 — 35 seats	11	General Track BoFs
Workshop Venue — 35–40 seats	8	The 8 pre-conference Workshops

Every one of the 75 devroom sessions — regardless of track, depth, or how oversubscribed its CFP was — runs in a room capped at 120 seats. The General Track, by contrast, splits between Hall 1 and Hall 2, averaging a 411-seat reach — more than 3.4× any devroom session. Weighting reach by duration ("seat-minutes" = reach × minutes) makes the gap concrete: the General Track alone accounts for roughly 78% of the conference's total seat-minutes of exposure, despite being only 40% of sessions, while the 8 devrooms — 54% of session count — account for under 19%. RTOS runs the deepest average content on the whole programme (7.4/10) but, like every other devroom, can never put that content in front of more than 120 people in a single sitting; a 30-minute General Track talk on any topic structurally outreaches it.

This has a direct bearing on the companion improvement plan's venue-math section: Hardware and Cloud & DevOps are independently identified there as this year's two most oversubscribed devrooms (an estimated ~75% rejection rate each), yet both are stuck at the same 120-seat ceiling as AOSP, this year's least oversubscribed devroom. A longer slot or a bigger room for either would raise reach and relieve oversubscription in the same move — see if2027-improvement-plan.pdf, Sections 2–3.

One session deserves a specific callout: **Communi-Con** (Section 8's forum-sourced CFP context aside, this is a General Track entry), a 75-minute community-voted lightning-talk slot whose CFP opens the day before the conference and closes on Day 1 — is booked into Hall 1, the conference's biggest room. It is, by seat-minutes, the single most-amplified slot on the entire 138-session programme, and as of this report its content is entirely undetermined: seven winning pitches will be picked by community vote after the main programme is already locked. That is either the most democratic slot on the schedule or the biggest reach-allocation wildcard, depending on how the vote turns out — worth watching.

Full Session Metrics Table

The complete 138-row table, including the new Reach column, follows on the next pages (landscape orientation).

Full Session Table — All 138 Scheduled Sessions

All 138 sessions (General Track, 8 devrooms, and the 8 pre-conference Workshops), classified against the categories.md schema plus this edition's new Reach column (seat capacity of the hall/room). Sorted by track, then time. Title links to the original CFP proposal. Multi-value fields separated by ",". Communi-Con's content is not yet determined (community vote after its CFP opens 25 Sep 2026).

Track	Session	Day	Min	Reach	Format	Platform	Domain	Stack Level	Category	Depth	Novelty	Audience
General Track	Can We Build a Copyleft License for Community Data in the Age of AI	day_2	60	35	Birds of Feather(BoF)	N/A	Legal/Policy;Research	N/A	FOSS Philosophy/Principles	5	7	researchers, business people, government employees
General Track	Community-Centred Technology	day_2	60	35	Birds of Feather(BoF)	embedded;mobile	Not-for-Profit Sector;Community	Application	Software	4	6	social workers, researchers
General Track	Do Open Source Companies Pay Competitively?	day_1	60	35	Birds of Feather(BoF)	N/A	Industry	N/A	FOSS Philosophy/Principles	3	4	software developers, business people
General Track	FOSS4Good - Techies in the social sector	day_1	60	35	Birds of Feather(BoF)	N/A	Not-for-Profit Sector;Community	N/A	FOSS Philosophy/Principles	3	4	social workers, software developers
General Track	Micro-granting in India	day_2	60	35	Birds of Feather(BoF)	N/A	Not-for-Profit Sector;Community	N/A	FOSS Philosophy/Principles	3	4	business people, social workers
General Track	How do we scale and sustain AI adoption without burning the planet?	day_2	60	35	Birds of Feather(BoF)	N/A	Research;Industry	N/A	FOSS Philosophy/Principles	4	6	researchers, business people, government employees
General Track	United by (Free) F&B – Food & Books	day_2	45	35	Birds of Feather(BoF)	N/A	Community	N/A	FOSS Philosophy/Principles	1	2	software developers, hackers
General Track	FOSS for the Forgotten Middle: Open-Sourcing Finance Software for India's 50,000 Microgyms	day_1	30	250	Talk	web;mobile	Industry;Community	Application	Software	5	6	business people, software developers
General Track	Navigating the FOSS Contribution Journey	day_1	30	250	Talk	N/A	Community	N/A	FOSS Philosophy/Principles	3	3	software developers, graduate/master level students
General Track	How FOSS and AI are Changing Government Workflows at NHAi	day_1	30	250	Talk	web	Government	Application	Software	5	6	government employees, software developers
General Track	Offline-first mobile apps on Frappe: a metadata-driven Flutter SDK with server-side control	day_1	30	250	Talk	mobile	Industry	Frameworks/Middleware	Software	7	6	software developers
General Track	Physics at 60FPS	day_1	30	250	Talk	desktop;web	Community;Industry	Libraries	Software	8	7	software developers, hackers
General Track	Calibrate: A framework for continuously evaluating and improving AI agents within non-profits	day_2	30	250	Talk	web	Not-for-Profit Sector;Industry	Application	Software	6	7	software developers, researchers
General Track	Being Bayesian: Stress-testing India's Loudest Stories with FOSS & Open Data	day_2	30	250	Talk	web	Research;Community	Application	Data	6	6	researchers, business people
General Track	Judiciary Entity Map: How open source can make India's courts legible	day_2	30	250	Talk	web	Government;Legal/Policy;Not-for-Profit Sector	Application	Data	6	7	researchers, government employees, business people
General Track	Software Patents - AI and the Future of Open Innovation	day_2	15	250	Lightning Talk	N/A	Legal/Policy	N/A	FOSS Philosophy/Principles	5	5	business people, government employees, software developers
General Track	Building Open Source Digital Infrastructure for a University	day_2	15	250	Lightning Talk	web;mobile	Education/Academia;Community	Application;Frameworks/Middleware	Software	6	5	software developers, graduate/master level students
General Track	Communi-Con : CFP opens 25th Sep	day_2	75	840	Lightning Talk	N/A	Community	N/A	N/A (content not yet determined)	—	—	N/A
General Track	From DPI to People's Digital Infrastructure: Reclaiming Digital Infrastructures for Democratic Accountability	day_1	60	840	Panel Discussion	N/A	Government;Legal/Policy;Community	N/A	FOSS Philosophy/Principles	4	5	government employees, social workers, researchers
General Track	The future of FOSS, SWE and Technical Education	day_2	60	840	Panel Discussion	N/A	Education/Academia;Industry	N/A	FOSS Philosophy/Principles	4	6	educators/professors, software developers, business people

General Track	FOSS and Co-operatives	day_1	60	35	Birds of Feather(BoF)	N/A	Legal/Policy;Community	N/A	FOSS Philosophy/Principles	4	5	business people, social workers
General Track	FOSS in the Public Consciousness	day_1	60	35	Birds of Feather(BoF)	N/A	Community	N/A	FOSS Philosophy/Principles	2	3	business people, social workers, software developers
General Track	Community-Driven or Community-Branded? Who Gets to Shape Community Events?	day_2	60	35	Birds of Feather(BoF)	N/A	Community	N/A	FOSS Philosophy/Principles	3	3	business people, software developers
General Track	FOSS for Justice — What can Open Source do for India's Courts?	day_2	60	35	Birds of Feather(BoF)	N/A	Government;Legal/Policy;Not-for-Profit Sector	N/A	FOSS Philosophy/Principles	4	5	government employees, social workers, software developers
General Track	Building OpenSpeaks tools for language documentation and archiving	day_1	30	840	Talk	web;desktop	Research;Community	Application	Software	5	6	researchers, software developers
General Track	Open Code, Curated Data: Building Open Water Intelligence for Indian Cities	day_1	30	840	Talk	web	Government;Community	Application	Data	6	6	government employees, researchers, software developers
General Track	Telling India's story through data	day_1	45	840	Invited Talk	web	Research;Community	Application	Data	5	5	researchers, business people, government employees
General Track	From Pull Requests to Patients: Building Open Healthcare Infrastructure in the Open	day_1	30	840	Talk	web	Not-for-Profit Sector;Government	Application	Software	6	5	software developers, social workers, government employees
General Track	From Institution to Community: Building COSH and a Culture of Open Source	day_1	30	840	Talk	N/A	Education/Academia;Community	N/A	FOSS Philosophy/Principles	3	3	educators/professors, graduate/master level students
General Track	Textbooks that run: interactive, zero-install CS education in the browser	day_1	30	840	Talk	web	Education/Academia	Application;Frameworks/Middleware	Software	6	6	educators/professors, software developers
General Track	Tracking-free and hackable network cameras - powered by libre software	day_1	30	840	Talk	embedded	Community;Industry	Firmware	Hardware	7	6	hackers, software developers
General Track	Fork Around and Find Out: How I Built a GNOME Extension Used by 200K+ People	day_1	15	840	Lightning Talk	desktop	Community	Application	Software	4	3	software developers, hackers
General Track	It takes a village to raise a repo	day_1	15	840	Lightning Talk	N/A	Community	N/A	FOSS Philosophy/Principles	3	3	software developers, social workers
General Track	Building Zasper: Inception to Growth	day_1	30	250	Talk	desktop	Community;Industry	Application	Software	7	6	software developers
General Track	Repub: Launching the Open Post-Processing Workflow for Digital Archiving in India	day_1	30	250	Talk	web	Not-for-Profit Sector;Research	Application	Software	5	6	researchers, software developers
General Track	How to Build Long-Lasting Tech Communities That Thrive — A Playbook	day_1	30	250	Talk	N/A	Community	N/A	FOSS Philosophy/Principles	3	3	business people, software developers
General Track	Optimizing Performance of GUIs	day_1	30	250	Talk	embedded;desktop	Industry	Libraries;Frameworks/Middleware	Software	7	6	software developers, hackers
General Track	How India's Largest G2C Platform Migrated to FOSS and Broke Free from Oracle	day_2	30	840	Talk	web	Government;Industry	Application;Frameworks/Middleware	Software	7	5	government employees, software developers, business people
General Track	How exactly can we teach students the right FOSS etiquette?	day_2	30	840	Talk	N/A	Education/Academia;Community	N/A	FOSS Philosophy/Principles	3	4	educators/professors, graduate/master level students
General Track	Seqout: search engine for the world's genomic data	day_2	30	840	Talk	web	Research	Application	Data	7	7	researchers, software developers
General Track	How we build Community Projects from the Community for the Community	day_2	30	840	Talk	web	Community	Application	Software	4	4	software developers, business people
General Track	Flexibility is complexity in disguise: Building frappe-ui for 10+ Products	day_2	30	840	Talk	web	Industry	Frameworks/Middleware;Libraries	Software	7	6	software developers

General Track	Surviving as an Open Source Startup : Xeneva	day_2	30	250	Talk	embedded;desktop	Industry	Kernel;Operating System	Software	6	5	business people, software developers, hackers
General Track	Computationalizing landscapes: The CoRE stack data structure for social-ecological diagnostics and planning	day_2	30	250	Talk	web	Research;Government;Not-for-Profit Sector	Application	Data	7	7	researchers, government employees
General Track	Freedom to code - Community's role in tech Policy	day_2	30	250	Talk	N/A	Legal/Policy;Community	N/A	FOSS Philosophy/Principles	6	5	business people, government employees, software developers
General Track	FQ-PIE: Contributing to Open Standards by Contributing to FOSS	day_1	15	250	Lightning Talk	N/A	Research;Industry	Kernel;Operating System	Software	8	7	researchers, software developers
General Track	Mapping India in seconds, for one and all	day_1	15	250	Lightning Talk	web	Research;Government	Application	Data	5	5	researchers, government employees
General Track	What can we do about open weight models enabling non consensual intimate imagery?	day_2	15	840	Lightning Talk	N/A	Research;Legal/Policy	Application	FOSS Philosophy/Principles	5	7	researchers, social workers, government employees
General Track	Opening the Archives: Building Digital Infrastructure for Public Participation in Archival Records	day_2	15	840	Lightning Talk	web	Not-for-Profit Sector;Research	Application	Data	5	6	researchers, software developers
General Track	Don't Model Every Leaf: The Magic of Blender Geometry Nodes	day_2	15	250	Lightning Talk	desktop	Community	Application	Software	5	4	designers, hackers
General Track	Property-based testing for UIs	day_2	15	250	Lightning Talk	web;mobile	Industry	Application;Libraries	Software	6	6	software developers
General Track	OpenTCAD: A FOSS Python Pipeline for Semiconductor Process and Device Simulation	day_2	15	250	Lightning Talk	desktop	Research;Industry	Application	Hardware	8	8	researchers, hackers
General Track	Open, decentralized language assessment software for the public school system	day_2	15	250	Lightning Talk	mobile;embedded	Education/Academia;Government	Application	Software	6	7	educators/professors, government employees
General Track	Bad Internet, Good Software	day_2	15	250	Lightning Talk	desktop	Community	Application	Software	5	4	hackers, software developers
General Track	Open Standards in XR - A current overview of OpenXR and WebXR	day_2	15	250	Lightning Talk	mobile;desktop;web	Industry;Community	Frameworks/Middleware	Software	5	4	software developers
AOSP	Launcher3 Internals: Understanding Android's Home Screen, Google Feed Integration, and Custom -1 Panels	day_1	25	120	Lightning Talk	mobile	Community;Industry	Application	Software	6	6	software developers, professionals
AOSP	Porting AOSP to Allwinner Sunxi platform	day_1	25	120	Talk	embedded;mobile	Community;Industry	Kernel;Operating System	Software	9	8	software developers, hackers, professionals
AOSP	Beyond HIDL: Modernizing Android HALs with Rust & Stable AIDL	day_1	25	120	Talk	embedded;mobile	Industry	Firmware;Kernel	Software	8	7	software developers, professionals
AOSP	Under the Hood of Android , What Really Happens When You Open an App?	day_1	15	120	Lightning Talk	mobile	Education/Academia;Community	Operating System	Software	5	4	software developers, graduate/master level students
AOSP	Inside Android Keyboards: Building an Open Source IME	day_1	15	120	Lightning Talk	mobile	Community	Application;Frameworks/Middleware	Software	6	6	software developers, hackers
AOSP	Adapting Generic Boot Loader on Android Platforms	day_1	15	120	Lightning Talk	embedded;mobile	Industry	Firmware;Operating System	Software	8	7	software developers, professionals
AOSP	Living With a DeGoogled Phone: Privacy, Trade-offs, and Daily Use	day_1	15	120	Lightning Talk	mobile	Community	Application	Software	3	3	hackers, social workers
AOSP	Volume Controller using Knob (Potentiometer)	day_1	15	120	Lightning Talk	embedded;mobile	Community;Education/Academia	Hardware;Firmware	Hardware	6	6	hackers, graduate/master level students
AOSP	Zero Trust, Zero TEE: Porting Android 16 to a Device Whose Security Stack Doesn't Exist Yet	day_1	15	120	Lightning Talk	embedded;mobile	Industry	Firmware;Operating System	Software	9	8	hackers, software developers, professionals
AOSP	Bypassing Android MTP: pushing a native C++ daemon via ADB for fast storage analysis	day_1	15	120	Lightning Talk	mobile;desktop	Community	Application	Software	6	6	hackers, software developers
AOSP	Keeping up with the AOSP	day_1	10	120	Lightning Talk	mobile	Community;Industry	Operating System	Software	4	3	software developers, professionals

Compilers	Bend, don't Break. Modern Techniques for Flexibility in Strongly Typed Functional Programming	day_1	20	120	Talk	N/A	Research;Community	Libraries;Frameworks/Middleware	Software	7	5	software developers, researchers
Compilers	PeerDAS: Erasure Coding When You Can't Trust the Storage	day_1	20	120	Talk	N/A	Research;Industry	Application;Libraries	Software	7	6	software developers, researchers
Compilers	Rust Is Not an HDL, So I Built a Tiny One.	day_1	20	120	Talk	N/A	Community;Research	Frameworks/Middleware;Hardware	Software	7	7	software developers, hackers
Compilers	Formal Verification: An Introduction through Elementary Arithmetic	day_1	20	120	Talk	N/A	Education/Academia;Research	Application	Software	6	5	researchers, educators/professors, graduate/master level students
Compilers	Semantic Quantization in KV Cache Eviction	day_1	15	120	Lightning Talk	N/A	Research;Industry	Application;Frameworks/Middleware	Software	7	7	software developers, researchers
Compilers	Databases in Ethereum Nodes	day_1	15	120	Lightning Talk	N/A	Research;Industry	Application	Software	7	6	software developers, researchers
Compilers	Benchmarking can be hard, here's why we're still creating a new one for OCaml	day_1	15	120	Lightning Talk	N/A	Research;Education/Academia	Libraries	Software	6	5	researchers, software developers
Compilers	What LLVM Does Differently When The Target Is CPU VS GPU	day_1	15	120	Lightning Talk	N/A	Research;Industry	Frameworks/Middleware;Kernel	Software	8	6	software developers, researchers
Compilers	PythonBPF - now with loops, custom funcs, and more maps!	day_1	15	120	Talk	N/A	Community;Industry	Kernel;Frameworks/Middleware	Software	7	7	software developers, hackers
Compilers	LGTM Is Not A Theorem	day_1	15	120	Lightning Talk	N/A	Industry	N/A	FOSS Philosophy/Principles	4	4	software developers, professionals
Compilers	One Language, All GPU: How Julia Targets different GPU backends from a Single Kernel.	day_1	15	120	Lightning Talk	N/A	Research;Community	Libraries;Frameworks/Middleware	Software	8	7	researchers, software developers
Documentation & Technical Writing	Documentation That Drives OSS Adoption: Lessons From Bruno	day_1	20	120	Talk	N/A	Industry;Community	N/A	FOSS Philosophy/Principles	4	4	software developers, business people
Documentation & Technical Writing	Developer Experience is the New Open Source Contribution. Here's Why	day_1	20	120	Talk	N/A	Community;Industry	N/A	FOSS Philosophy/Principles	3	4	software developers
Documentation & Technical Writing	Writing Docs for Two Readers: Humans and AI Agents	day_1	20	120	Talk	N/A	Industry;Community	Application	Software	5	7	software developers
Documentation & Technical Writing	404: Empathy Not Found	day_1	20	120	Talk	N/A	Community;Industry	N/A	FOSS Philosophy/Principles	3	4	software developers
Documentation & Technical Writing	Your first open source contribution does not have to be code	day_1	10	120	Lightning Talk	N/A	Community	N/A	FOSS Philosophy/Principles	2	3	graduate/master level students, software developers
Documentation & Technical Writing	Why Documentation Shouldn't Feel Like Plain Text	day_1	10	120	Lightning Talk	N/A	Community;Industry	Application	Software	5	6	software developers
Documentation & Technical Writing	Ctrl+F for 34 Years of India's GovTech (Building a Public Commons for India's Digital Governance History)	day_1	10	120	Lightning Talk	web	Government;Research	Application	Data	5	6	government employees, researchers
Documentation & Technical Writing	Refactoring documentation without breaking it	day_1	10	120	Lightning Talk	N/A	Community;Industry	Application	Software	5	5	software developers
Documentation & Technical Writing	Changelogs Are Not Enough: Writing Community Memory for Kubernetes	day_1	10	120	Lightning Talk	N/A	Community;Industry	N/A	FOSS Philosophy/Principles	4	5	software developers
Documentation & Technical Writing	Python Now Speaks Tamil: Getting Official Docs Translated for a Mother Tongue	day_1	10	120	Lightning Talk	N/A	Education/Academia;Community	N/A	FOSS Philosophy/Principles	3	5	educators/professors, graduate/master level students
Documentation & Technical Writing	Printable, Downloadable, Forever: An Ode to Offline Documentation	day_1	10	120	Lightning Talk	N/A	Community	N/A	FOSS Philosophy/Principles	2	3	software developers
Cloud & DevOps	Remote Execution by Design: Sandboxing AI Agents with gVisor & Pipelock	day_2	30	120	Talk	N/A	Industry	Kernel;Operating System	Software	8	7	software developers, professionals

Cloud & DevOps	How OpenTelemetry Instrumentation Works Under the Hood	day_2	15	120	Lightning Talk	N/A	Industry	Frameworks/Middleware;Libraries	Software	7	6	software developers
Cloud & DevOps	From Prompts to Production: Building an Open-Source AI On-Call Agent with MCP and Multi-Agent Workflows	day_2	15	120	Lightning Talk	N/A	Industry	Application;Frameworks/Middleware	Software	6	6	software developers
Cloud & DevOps	Long live outbox : Outbox is dead	day_2	15	120	Lightning Talk	N/A	Industry	Application	Software	8	7	software developers, professionals
Cloud & DevOps	Kubernetes Native Observability Dashboards with Perses, without the YAML/JSON Hell	day_2	15	120	Lightning Talk	N/A	Industry	Application;Frameworks/Middleware	Software	6	5	software developers
Cloud & DevOps	Cilium Network Policy Anti-Patterns I Learned the Hard Way	day_2	15	120	Lightning Talk	N/A	Industry	Kernel;Operating System	Software	7	5	software developers, professionals
Cloud & DevOps	Escaping Cron Hell: Building and Contributing to Dagu, an Open Source Workflow Orchestrator	day_2	15	120	Lightning Talk	N/A	Industry;Community	Application	Software	6	5	software developers
Cloud & DevOps	15 Years, One Server, 8GB RAM and 500,000 Linux Learners: Lessons from Running Webminal	day_2	15	120	Lightning Talk	web	Education/Academia;Community	Operating System;Application	Software	5	4	educators/professors, software developers
Cloud & DevOps	Unused Is Not a Kubernetes Field: Platform Hygiene Is Not Cleanup	day_2	15	120	Lightning Talk	N/A	Industry	Application	Software	6	5	software developers, professionals
Cloud & DevOps	I Deleted Production With helm upgrade. Here's the Exact Chain of Events.	day_2	15	120	Lightning Talk	N/A	Industry	Application	Software	5	4	software developers, professionals
RTOS	From Reset to main(): Porting Infineon TriCore AURIX to Zephyr	day_2	25	120	Talk	embedded	Industry	Kernel;Firmware	Software	9	8	software developers, professionals
RTOS	Beyond Hardware Description: Using Device Tree for System-Level Configuration in Zephyr RTOS	day_2	15	120	Lightning Talk	embedded	Industry	Firmware;Operating System	Software	7	6	software developers, professionals
RTOS	From Display Controller to Panel: A Modular Display Pipeline Architecture for Zephyr	day_2	15	120	Lightning Talk	embedded	Industry	Firmware;Kernel	Software	7	6	software developers, professionals
RTOS	Real-Time Video Pipelines on Embedded SoCs: A FreeRTOS Case Study	day_2	15	120	Lightning Talk	embedded	Industry	Firmware;Kernel	Software	8	6	software developers, professionals
RTOS	Tackling core lifecycle management in heterogeneous SoCs	day_2	15	120	Lightning Talk	embedded	Industry	Kernel;Firmware	Software	8	7	software developers, professionals
RTOS	Loops to real time	day_2	15	120	Lightning Talk	embedded	Education/Academia;Community	Firmware	Software	4	3	hackers, graduate/master level students
RTOS	Indigenous, High-Performance, Safety-Critical RTOS	day_2	15	120	Lightning Talk	embedded	Government;Industry	Kernel;Operating System	Software	8	8	government employees, software developers, professionals
RTOS	From Architecture to Tasks: Designing Real-Time MCU Firmware in Heterogeneous APU-MCU Systems	day_2	15	120	Lightning Talk	embedded	Industry	Firmware;Kernel	Software	7	6	software developers, professionals
RTOS	Deterministic Vision: Offloading Real-Time Camera Pipelines to FreeRTOS on Heterogeneous SoCs	day_2	15	120	Lightning Talk	embedded	Industry	Firmware;Kernel	Software	8	7	software developers, professionals
RTOS	RTL to Zephyr RTOS: Lightweight-Crypto Acceleration on Open RISC-V	day_2	15	120	Talk	embedded	Research;Industry	Hardware;Firmware	Hardware	8	8	software developers, researchers, hackers
Security	Keeping an Eye on your Eye: Disassembling and Reassembling Smart Camera Hardware to breach YOUR Security	day_2	30	120	Talk	embedded	Community	Hardware;Firmware	Hardware	7	6	hackers, software developers
Security	The Art of Userspace Sandboxing on Linux	day_2	25	120	Talk	desktop	Industry;Community	Kernel;Operating System	Software	8	6	software developers, hackers, professionals

Security	Papers, Please: A First Look at Age Verification on the Web	day_2	25	120	Talk	web	Research;Legal/Policy	Application	Data	7	7	researchers, government employees, business people
Security	Extracting AES Keys from Embedded Devices Using Open Source Hardware	day_2	20	120	Lightning Talk	embedded	Research;Education/Academia	Hardware;Firmware	Hardware	8	7	hackers, researchers, graduate/master level students
Security	Reverse Engineering the PAN QR Code	day_2	15	120	Lightning Talk	mobile;desktop	Government;Community	Application	Software	6	7	hackers, software developers
Security	When NO_PROXY Lies: Finding an IPv4-Mapped IPv6 Patch Bypass in Axios	day_2	15	120	Lightning Talk	web	Industry;Community	Libraries	Software	7	7	software developers, hackers
Security	Extending Capa for Malicious Script Analysis	day_2	15	120	Lightning Talk	desktop	Industry;Community	Application;Libraries	Software	7	6	software developers, hackers
Security	Hacking India's Largest Exam System	day_2	15	120	Lightning Talk	web	Government;Community	Application	Software	6	6	hackers, government employees, software developers
Security	A plan to sort out this CVE mess	day_2	15	120	Lightning Talk	N/A	Community;Industry	N/A	FOSS Philosophy/Principles	5	6	software developers, business people
Design	Forkable design	day_2	20	120	Talk	N/A	Industry;Community	Application	Software	5	5	designers, software developers
Design	Shaili Sanchaya: Reviving Indic Typography from Archival Metadata	day_2	20	120	Talk	web	Research;Community	Application	Data	6	7	designers, researchers
Design	Designing for Reality: Where UI Meets the Physical World	day_2	20	120	Talk	embedded;mobile	Industry	Application	Software	5	5	designers, software developers
Design	Creating Design Systems in p5JS to accelerate Content Creation	day_2	20	120	Talk	web	Industry;Community	Application;Libraries	Software	5	5	designers, business people
Design	KaiOS Design Notes: UX observations from a feature phone	day_2	20	120	Talk	mobile	Community	Application	Software	5	6	designers, software developers
Design	Making Desktops Expressive: Introducing Bézier Curves to GTK	day_2	20	120	Talk	desktop	Community	Libraries;Application	Software	6	6	designers, software developers, hackers
Design	Sound Is All You Need: From Good UX to Great UX	day_2	20	120	Talk	N/A	Industry	Application	Software	4	5	designers
Hardware	Build your own open source keyboard with ZMK/QMK	day_1	30	120	Talk	embedded	Community	Firmware;Hardware	Hardware	5	4	hackers
Hardware	Building Open Source Music Hardware in India - Tarab Instruments	day_1	30	120	Talk	embedded	Community;Industry	Hardware;Firmware	Hardware	6	6	hackers, business people
Hardware	CNC4Everyone - How we built a CNC machine for every classroom, college and makerspace	day_1	30	120	Talk	embedded	Education/Academia;Community	Hardware;Firmware	Hardware	6	6	hackers, educators/professors, graduate/master level students
Hardware	My zero-to-hero journey deploying a fiber-optic and wireless community mesh	day_1	30	120	Talk	embedded	Community;Not-for-Profit Sector	Hardware;Firmware	Hardware	5	5	hackers, social workers
Hardware	HackerFab IITB: democratizing semiconductor fabrication	day_1	30	120	Talk	embedded	Education/Academia;Research	Hardware	Hardware	8	8	researchers, graduate/master level students, hackers
Hardware	Minnow: An Open Platform for Ground Swarm Robotics	day_1	25	120	Talk	embedded	Research;Education/Academia	Hardware;Application	Hardware	6	7	hackers, researchers, graduate/master level students
Workshops (Pre-Conference Day)	Getting Started with Qt	day_0	120	38	Workshop	desktop	Community;Industry	Frameworks/Middleware	Software	5	3	software developers, hackers
Workshops (Pre-Conference Day)	Making music with strudel.cc	day_0	120	38	Workshop	web	Community;Education/Academia	Application	Software	4	6	hackers, designers, graduate/master level students
Workshops (Pre-Conference Day)	Digital Security for Journalists, Activists, and Everyday Internet Users: A Hands-On FOSS Workshop	day_0	120	38	Workshop	mobile;desktop	Not-for-Profit Sector;Community	Application	Software	5	4	social workers, hackers

Workshops (Pre-Conference Day)	Fun and Profit with OCaml	day_0	120	38	Workshop	desktop	Education/Academia;Research	Application;Libraries	Software	6	5	software developers, graduate/master level students
Workshops (Pre-Conference Day)	Art - meet Math - meet Code	day_0	120	38	Workshop	web	Education/Academia;Community	Application	Software	4	5	designers, hackers, educators/professors
Workshops (Pre-Conference Day)	Building AI agents end-to-end: A hands-on workshop	day_0	120	38	Workshop	web	Industry	Application;Frameworks/Middleware	Software	6	6	software developers
Workshops (Pre-Conference Day)	Adding places to OpenStreetMap (outdoor workshop)	day_0	120	38	Workshop	mobile	Community	Application	Data	3	3	hackers, social workers
Workshops (Pre-Conference Day)	Self-Hosting 101: A real-time self-hosting workshop	day_0	120	38	Workshop	desktop	Community	Application;Operating System	Software	5	3	hackers, software developers

Note on Methodology & Categorization

This report was produced by an AI assistant (Claude, by Anthropic) working through the conference's own source files, not by the IndiaFOSS 2026 programme committee. It is intended as an analytical starting point for programme review, not as an authoritative or final judgement on any session, speaker, or track.

How the session list was built

1. **Schedule extraction.** if2026-schedule.ods was parsed sheet-by-sheet (General Track plus the 8 devroom sheets, excluding the Food Breaks sheet). Day, hall and start-time values are only written on the first row of each schedule block in the source file, so they were forward-filled down subsequent rows within a block; the per-row exclude flag (General Track only) was applied per-row, not forward-filled, to avoid dropping unrelated sessions that merely followed an excluded one.
2. **Title matching.** Each session title was normalised (smart quotes, dashes and whitespace collapsed) and matched against cfp-info.csv to recover its CFP ID. All 130 main-programme sessions and all 8 pre-conference workshops matched an ID; nothing was matched by fuzzy/approximate title similarity.
3. **Proposal retrieval.** For every matched ID, the full public proposal page at fossunited.org/c/indiafoss/2026/cfp/<ID> was fetched and its "Session Description" field was extracted as plain text. All scoring and tagging below was based on reading this original proposal text, not on the title alone. One exception: the "Communi-Con : CFP opens 25th Sep" entry has a confirmed slot, duration and room, but no describable content — its own CFP doesn't open until the day before the conference and the seven winning pitches are chosen by community vote — so it is excluded from every category/depth/novelty breakdown in this report and shown as "not yet determined" in the full table.
4. **Workshops.** The 8 Workshop-type entries in cfp-info.csv do not appear in if2026-schedule.ods because they run on a separate pre-conference day; they were added to the dataset manually as a 10th track ("Workshops (Pre-Conference Day)") with an assumed uniform duration of 120 minutes each, per the stated ~2-hour format, since no individual start times were published for them.
5. **Reach.** New this edition: each session's room (from the schedule) was mapped to that room's published seat capacity — Hall 1: 840, Hall 2: 250, Hall 3 & Room 1: 120, Room 2 & 3: 35 — and the 8 pre-conference Workshops, which run in a separate dedicated venue, were assigned a capacity of 35–40 seats per the stated workshop-venue sizing. This is a room-capacity figure, not a verified attendance count — it says how many people *could* be in the room, not how many were.
6. **Curation attribution.** The report tags each gap/strength as "Programme Committee" (General Track, BoFs, Panels and the Workshop day) or "Independent devroom" (the 8 devrooms, each running its own internal review once selected). The named rejected proposals (Section 8 of the Detailed Overview) come from three public FOSS United forum threads: the call-for-devrooms-for-indiafoss-2026, propose-birds-of-a-feather-bof-session-indiafoss-2026, and call-for-panel-discussions-indiafoss-2026 threads. Those threads were read via an automated fetch-and-summarise tool, not parsed directly from the raw HTML, so treat the exact proposal counts and titles as a close paraphrase of the forum record rather than a verbatim transcript.
7. **Booths and sponsors (Section 9).** The 75-booth list at fossunited.org/indiafoss/2026/booths and the sponsor/community-partner logos at fossunited.org/indiafoss/2026 are treated separately from the session dataset throughout this report — neither booths nor sponsors are talks, so the categories.md schema (duration, depth, novelty, reach) doesn't apply to them.
8. **Which devrooms are community-run vs. individual-run.** The distinction drawn in Section 1 and the gap analysis was supplied as fact for this analysis rather than independently sourced by this report — the fossunited.org/indiafoss/2026/devrooms index page itself doesn't name managers or communities per devroom, so this detail couldn't be independently cross-checked from public pages alone.
9. **Speaker counts and the producer-to-attendee ratio (Section 10).** The 157-unique-speaker figure was extracted directly from the "Speakers" field already present on each of the 138 cached session proposal pages. The attendee statistics (2,401 total registrations, ticket tiers, gender and regional breakdowns) come from fossunited.org/indiafoss/2026/stats, re-fetched for this edition after the project's own documentation clarified that the gap between the ticket-tier sum and the headline total is 760 reserved free tickets for speakers, volunteers, sponsors, booth staff and diversity scholars — confirmed rather than inferred, unlike in the previous edition. The 150-booth-representative figure is a nominal ceiling (75 booths × 2 seats), not a verified headcount, and cannot be de-duplicated against the speaker list.
10. **The BoF acceptance rate (Section 8).** The 14-submission, 3-rejection BoF figure comes from a public list of rejected CFP submissions (title, session type, status) provided directly as a file, the same way cfp-info.csv was.
11. **Speaker background buckets (Section 11).** Sourced from the same "Speakers" field used for the Section 10 count, here also read for each speaker's stated designation and organisation. Sorting 157 one-line self-descriptions into 9 buckets (student, academic/researcher, large/established company, startup/SME employee, founder/entrepreneur, non-profit/civil-society, government, independent/freelance, not stated/unclear) is a single analyst's judgement call, not an HR-verified classification. This edition applied one explicit rule to resolve the founder-vs-employee ambiguity flagged in earlier editions:

the founder/entrepreneur bucket is used only where the stated title itself says "Founder" or "Co-Founder" — a "CEO" or "Founding Engineer" without that word is filed under startup/SME employee (or large/established company, if the organisation is one) instead. No attempt was made to infer gender, region, or anything not stated in that one line.

How each session was classified

Every one of the 138 sessions was individually read and tagged by the AI assistant against the seven dimensions defined in `categories.md`, plus this edition's new reach figure:

- **Platform, Domain, Level in Technology Stack** — assigned as one or more free-text tags per session (separated by ";" in the data table), based on the technology and audience the proposal itself describes. A session can and often does carry more than one tag in each field.
- **Overall category** — a single primary label (Hardware / Software / Data / FOSS Philosophy & Principles) chosen to reflect the proposal's centre of gravity, not every topic it touches.
- **Depth of technical coverage (1–10) and Novelty vs. existing work (1–10)** — judgement calls made by the AI assistant from the proposal text alone. These were not derived from any objective external signal (e.g. citation counts, GitHub stars, speaker seniority) and should be read as one analyst's relative ranking across the programme, not an absolute or verified score.
- **Suitability for audience** — one or more audience tags from the `categories.md` list, inferred from who the proposal explicitly addresses or clearly assumes as background.
- **Reach** — the seat capacity of the session's assigned room, read directly off the schedule and the conference's published room-capacity list; not a judgement call.

WHERE THIS ANALYSIS IS MOST LIKELY TO BE WRONG

- Depth and novelty scores are subjective single-pass estimates from proposal text, not from watching the actual delivered talk — a strong speaker can make a "5/10 novelty" topic land as fresh, and a weak delivery can flatten a "9/10" one.
- Classification was done from the CFP abstract only; last-minute scope changes, co-speakers, or slide content that diverges from the original proposal are not reflected.
- Multi-tag fields (platform, domain, stack level, audience) reflect the assistant's reading of emphasis, which is inherently approximate for sessions that genuinely span several areas.
- The dataset reflects a single snapshot of `if2026-schedule.ods` and `cfp-info.csv`; if the schedule changes after this report was generated, the numbers here will no longer match the live programme. This edition already reflects one such change from the previous edition — one new General Track entry (Communi-Con) was added to the schedule since the last snapshot.
- Reach reflects room capacity, not actual attendance or demand — a 120-seat devroom session that's consistently full is, in a real sense, "reaching" its ceiling just as hard as an under-attended 840-seat General Track talk; capacity alone doesn't distinguish the two.
- The speaker-bucket classification in Section 11 was fully re-done for this edition against the current speaker data rather than carried over from the previous edition, because several designation fields that were reportedly blank in the prior pass now carry real text — plausibly because CFP speaker profiles are editable and were updated between editions. The bucket counts here should be treated as this edition's own numbers, not identical to any earlier edition's, since the classification is a fresh single-analyst read of possibly-changed source text.

The full per-session table earlier in this report is provided precisely so that any individual classification can be checked against the linked original proposal and corrected by someone with direct knowledge of the talk.

Appendix: categories.md, verbatim, and how it maps to this report

Every classification in this report traces back to a single project file, categories.md, reproduced below exactly as written (including its two small typos — "Not-for-Profot" and "FOSS Philosphy/Principles" — left uncorrected because this is a verbatim quote, not a paraphrase).

CATEGORIES.MD — FULL TEXT

Use the following categories for classification:

- Platform: embedded, mobile, desktop, web, etc
- Domain: Education/Academia, Not-for-Profot Sector, Government, Research, Community, Industry, Legal/Policy, etc
- Level in Technology Stack: Hardware, Firmware, Kernel, Operating System, Frameworks/Middleware, Libraries, Application, etc
- Overall category: Hardware, Software, Data, FOSS Philosphy/Principles, etc
- Depth of technical coverage at Level: rated on a scale of 1 to 10
- Novelty compared to existing work: rated on a scale of 1 to 10
- Suitability for Audience: graduate/master level students, designers, hackers, professionals, researchers, software developers, educators/professors, social workers, government employees, venture capitalists, business people, etc

Each of the seven bullets became one column in the per-session table earlier in this report; reach is this report's own addition and isn't part of categories.md. The table below shows how the free-form category definitions were operationalised, and what this programme's 138 sessions show under that rule:

categories.md bullet	Report column	What this programme's 138 sessions show
Platform	platform	Web 32 · Embedded 29 · Mobile 23 · Desktop 19 · N/A (process/community sessions with no specific platform) 55 (tag instances; a session can carry >1 tag, and many BoFs/policy/community talks carry none).
Domain	domain	Community 76 · Industry 65 · Research 30 · Education/Academia 20 · Government 14 · Not-for-Profit Sector 12 · Legal/Policy 9.
Level in Technology Stack	stack_level	Application 65 · Firmware 20 · Frameworks/Middleware 19 · Kernel 16 · Operating System 14 · Libraries 14 · Hardware 11 (tag instances; 29 sessions carry no stack-level tag at all, being pure BoF/policy/community content).
Overall category	overall_category	Software 84 sessions / 2,370 min · FOSS Philosphy/Principles 29 / 1,030 min · Data 12 / 400 min · Hardware 12 / 300 min · not-yet-determined (Communi-Con) 1 / 75 min. One value was still chosen per session (not multiple) — this report's own reading of "overall category" as a centre-of-gravity label.
Depth of technical coverage	depth	Mean 5.6/10 across the 137 scoreable sessions; highest by track: RTOS 7.4, Security 6.8; lowest: Documentation & Technical Writing 3.7 and the Workshop day 4.8 (both accessible by design).
Novelty vs. existing work	novelty	Mean 5.5/10; highest by track: RTOS 6.5; lowest: the Workshop day 4.4, intentionally introductory.
Suitability for Audience	audience	software developers 99 · hackers 39 · researchers 33 · business people 23 · professionals 21 · government employees 20 · graduate/master level students 16 · social workers 14 · educators/professors 10 · designers 10 · venture capitalists 0 (tag instances; a session can carry >1 tag).
(not in categories.md — this report's own addition)	reach	Room 2 & 3 (35 seats) 11 sessions · Workshop Venue (35–40) 8 · Hall 3 / Room 1 (120) 75 · Hall 2 (250) 25 · Hall 1 (840) 19. Every devroom session, without exception, runs at the 120-seat tier.

All seven categories.md bullets carry an "etc." or an open numeric scale — there is no closed, exhaustive list in the source file. This report read "etc." on Platform, Domain, Level in Technology Stack, and Suitability for Audience as permission to tag more than one value per session and to extend the named examples where warranted. Overall category is the one exception by choice, not by rule: even though a 5th value is explicitly allowed, this report kept it to one of the four named values per session (plus "not yet determined" for the one session with no content to classify), because no session needed a genuinely new value and a single centre-of-gravity label stayed the most useful reading of what "overall category" means.